

Programming Assignment 4

Name: Muzammilkhon Muradullaev

```
In [22]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [23]: import os
import numpy as np
import torch

from torch.utils.data import DataLoader
from torchvision import datasets, transforms, models
from torchvision.models.feature_extraction import create_feature_extractor
```

```
In [24]: DATA_DIR = "/content/drive/MyDrive/Faulty_solar_panel"
print(os.listdir(DATA_DIR))
```

['test.json', 'train.json', 'validation.json', '.DS_Store', 'Clean', 'Snow-Covered', 'Dusty', 'venv', 'Bird-drop']

Dataset used (same as previous assignments)

I use the Faulty_solar_panel image dataset located at:

```
/content/drive/MyDrive/Faulty_solar_panel
```

It contains 4 classes (*Bird-drop*, *Clean*, *Dusty*, *Snow-Covered*).

This satisfies the assignment requirement to use the same dataset as previous assignments and matches K = 4 for clustering.

```
In [25]: ALLOWED_EXT = {".jpg", ".jpeg", ".png", ".bmp"}

def is_valid_image(path: str) -> bool:
    path = path.lower()
    ext = os.path.splitext(path)[1]
    return ext in ALLOWED_EXT
```

```
In [26]: transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]
    )
])
```

```
In [27]: dataset = datasets.ImageFolder(
    root=DATA_DIR,
    transform=transform,
    is_valid_file=is_valid_image
)
```

```

print("Classes:", dataset.classes)
print("Total images:", len(dataset))
print("First file:", dataset.samples[0][0])

Classes: ['Bird-drop', 'Clean', 'Dusty', 'Snow-Covered', 'venv']
Total images: 778
First file: /content/drive/MyDrive/Faulty_solar_panel/Bird-drop/Bird (1).jpeg

```

In [28]: `loader = DataLoader(dataset, batch_size=16, shuffle=False)`

In [29]: `device = "cuda" if torch.cuda.is_available() else "cpu"`

```

model = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
model.eval()

extractor = create_feature_extractor(
    model,
    return_nodes={"layer4": "features"}
).to(device)

```

In [30]: `@torch.no_grad()
def extract_features(loader):
 feats = []
 for imgs, _ in loader:
 imgs = imgs.to(device)
 out = extractor(imgs)["features"] # [B, 512, 7, 7]
 out = out.mean(dim=(2, 3)) # GAP → [B, 512]
 feats.append(out.cpu().numpy())
 return np.vstack(feats)

X = extract_features(loader)
print("Feature matrix shape:", X.shape)`

Feature matrix shape: (778, 512)

Feature extraction reference (required)

Feature extraction from ResNet18 follows the approach described here (required reference):

<https://kozodoi.me/blog/20210527/extracting-features>

I extracted activations from `layer4` (last convolutional block) and apply global average pooling to obtain a 512-D feature vector per image.

In [31]: `import numpy as np
y_true = np.array(dataset.targets)
print("y_true shape:", y_true.shape, "unique:", np.unique(y_true))`

y_true shape: (778,) unique: [0 1 2 3 4]

In [32]: `from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA`

```

X_scaled = StandardScaler().fit_transform(X)
X_2d = PCA(n_components=2, random_state=42).fit_transform(X_scaled)

```

```
print("X_2d shape:", X_2d.shape)
```

X_2d shape: (778, 2)

```
In [33]: import pandas as pd
from sklearn.metrics import fowlkes_mallows_score, silhouette_score
from sklearn.cluster import (
    KMeans, BisectingKMeans, SpectralClustering, DBSCAN, AgglomerativeClu
)

results = []

def safe_silhouette(X2d, labels):
    uniq = set(labels)
    if len(uniq) < 2:
        return np.nan
    try:
        return silhouette_score(X2d, labels)
    except Exception:
        return np.nan

def evaluate(name, labels):
    return {
        "Method": name,
        "FMI": fowlkes_mallows_score(y_true, labels),
        "Silhouette": safe_silhouette(X_2d, labels),
        "Clusters": len(set(labels))
    }
```

```
In [34]: labels = KMeans(n_clusters=4, init="random", n_init=10, random_state=42).fit(X_2d)
results.append(evaluate("KMeans (init=random)", labels))
```

```
In [35]: labels = KMeans(n_clusters=4, init="k-means++", n_init=10, random_state=42).fit(X_2d)
results.append(evaluate("KMeans (init=k-means++)", labels))
```

```
In [36]: labels = BisectingKMeans(n_clusters=4, init="random", random_state=42).fit(X_2d)
results.append(evaluate("BisectingKMeans (init=random)", labels))
```

```
In [37]: labels = SpectralClustering(n_clusters=4).fit_predict(X_2d)
results.append(evaluate("SpectralClustering (default params)", labels))
```

```
In [38]: def dbscan_find_params(X2d, target_k=4):
    eps_grid = np.linspace(0.1, 5.0, 60)
    min_samples_grid = [3, 4, 5, 6, 8, 10]

    best = None # (sil, eps, ms, labels, k)
    for ms in min_samples_grid:
        for eps in eps_grid:
            model = DBSCAN(eps=float(eps), min_samples=int(ms))
            labels = model.fit_predict(X2d)
            k = len(set(labels)) - (1 if -1 in labels else 0) # without noise

            if k == target_k:

                mask = labels != -1
                if mask.sum() < 2 or len(set(labels[mask])) < 2:
                    sil = -np.inf
                else:
```

```

try:
    sil = silhouette_score(X_2d[mask], labels[mask])
except Exception:
    sil = -np.inf

cand = (sil, eps, ms, labels, k)
if (best is None) or (cand[0] > best[0]):
    best = cand

return best

best_db = dbscan_find_params(X_2d, target_k=4)

if best_db is None:
    print("DBSCAN: cannot find the parameters which gives 4 clusters.")
else:
    sil, eps_used, ms_used, labels, k = best_db
    print(f"DBSCAN cannot find: eps={eps_used:.3f}, min_samples={ms_used}")

mask = labels != -1
results.append({
    "Method": f"DBSCAN (eps={eps_used:.3f}, min_samples={ms_used})",
    "FMI": fowlkes_mallows_score(y_true[mask], labels[mask]),
    "Silhouette": safe_silhouette(X_2d[mask], labels[mask]),
    "Clusters": k
})

```

DBSCAN cannot find: eps=0.266, min_samples=6, clusters=4

DBSCAN note (eps/min_samples + evaluation)

For DBSCAN, we search `eps` and `min_samples` to obtain exactly 4 clusters (excluding noise).

DBSCAN may label some points as noise (`-1`). For a fair comparison, FMI and Silhouette are computed on non-noise points only (`labels != -1`).

```
In [39]: for link in ["single", "complete", "average", "ward"]:
    labels = AgglomerativeClustering(n_clusters=4, linkage=link).fit_pred
    results.append(evaluate(f"Agglomerative (linkage={link})", labels))
```

```
In [40]: df = pd.DataFrame(results)

print("== ALL RESULTS ==")
display(df)

print("== Rank by FMI (best -> worst) ==")
display(df.sort_values("FMI", ascending=False).reset_index(drop=True))

print("== Rank by Silhouette (best -> worst) ==")
display(df.sort_values("Silhouette", ascending=False).reset_index(drop=True))
```

== ALL RESULTS ==

	Method	FMI	Silhouette	Clusters
0	KMeans (init=random)	0.392029	0.391367	4
1	KMeans (init=k-means++)	0.392029	0.391367	4
2	BisectingKMeans (init=random)	0.376632	0.359498	4
3	SpectralClustering (default params)	0.506796	0.473370	4
4	DBSCAN (eps=0.266, min_samples=6)	0.446203	0.934064	4
5	Agglomerative (linkage=single)	0.473063	0.393752	4
6	Agglomerative (linkage=complete)	0.393774	0.325200	4
7	Agglomerative (linkage=average)	0.491869	0.428786	4
8	Agglomerative (linkage=ward)	0.346901	0.320865	4



==== Rank by FMI (best -> worst) ===

	Method	FMI	Silhouette	Clusters
0	SpectralClustering (default params)	0.506796	0.473370	4
1	Agglomerative (linkage=average)	0.491869	0.428786	4
2	Agglomerative (linkage=single)	0.473063	0.393752	4
3	DBSCAN (eps=0.266, min_samples=6)	0.446203	0.934064	4
4	Agglomerative (linkage=complete)	0.393774	0.325200	4
5	KMeans (init=random)	0.392029	0.391367	4
6	KMeans (init=k-means++)	0.392029	0.391367	4
7	BisectingKMeans (init=random)	0.376632	0.359498	4
8	Agglomerative (linkage=ward)	0.346901	0.320865	4


==== Rank by Silhouette (best -> worst) ===

	Method	FMI	Silhouette	Clusters
0	DBSCAN (eps=0.266, min_samples=6)	0.446203	0.934064	4
1	SpectralClustering (default params)	0.506796	0.473370	4
2	Agglomerative (linkage=average)	0.491869	0.428786	4
3	Agglomerative (linkage=single)	0.473063	0.393752	4
4	KMeans (init=random)	0.392029	0.391367	4
5	KMeans (init=k-means++)	0.392029	0.391367	4
6	BisectingKMeans (init=random)	0.376632	0.359498	4
7	Agglomerative (linkage=complete)	0.393774	0.325200	4
8	Agglomerative (linkage=ward)	0.346901	0.320865	4


Distributions




2-d distributions




Time series





Values






WARNING: Runtime no longer has a reference to this dataframe, please re-run this cell and try again.


Distributions




2-d distributions




Time series




Values




Distributions




2-d distributions






Time series



Values



In [40]: